

# Transformers

Greg Strabel

October 6, 2023

## 1 Preliminaries

### 1.1 Matrix Multiplication

Given two matrices  $X \in \mathbb{R}^{n \times m}$  and  $Y \in \mathbb{R}^{m \times k}$

$$(XY)_{ij} = \sum_{l=1}^m X_{il}Y_{lj} = X_{i.}Y_{.j} \quad (1)$$

Therefore

$$(XY)_{.j} = \sum_{l=1}^m Y_{lj}X_{.l} \quad (2)$$

so that the columns of  $XY$  are linear combinations of the columns of  $X$  and

$$(XY)_{i.} = \sum_{l=1}^m X_{il}Y_{l.} \quad (3)$$

so that the rows of  $XY$  are linear combinations of the rows of  $Y$ .

### 1.2 Softmax

**Definition 1.1** (Softmax). The softmax function  $\sigma : \mathbb{R}^K \rightarrow \mathbb{R}^K$  is

$$\sigma(z)_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}} \quad (4)$$

### 1.3 Evaluating Text Quality

**Definition 1.2** (n-grams). The n-grams of a string  $y = y_1 \dots y_K$  are

$$G_n(y) = \{y_1 \dots y_n, y_2 \dots y_{n+1}, \dots, y_{K-n+1} \dots y_K\} \quad (5)$$

Note that  $G_n(y)$  is a set, not a multiset, so the elements of  $G_n(y)$  are distinct.

**Definition 1.3** (Substring Count). Given two strings,  $s = s_1 \dots s_M$  and  $y = y_1 \dots y_K$  the substring count of substring  $s$  in  $y$  is

$$C(s, y) = |\{i \in \mathbb{N} \mid i > 0, i + M \leq K, y_i \dots y_{i+M} = s\}| \quad (6)$$

**Definition 1.4** (Modified n-gram precision). Given a candidate corpus  $\hat{S} = (\hat{y}^{(1)}, \hat{y}^{(2)}, \dots, \hat{y}^{(M)})$  and corresponding reference corpi  $S = (S_1, S_2, \dots, S_M)$  where  $S_i = (y^{(i,1)}, y^{(i,2)}, \dots, y^{(i,N_i)})$ , modified n-gram precision is defined as

$$p_n(\hat{S}, S) = \frac{\sum_{i=1}^M \sum_{s \in G_n(\hat{y}^{(i)})} \min(C(s, \hat{y}^{(i)}), \max_{y \in S_i} C(s, y))}{\sum_{i=1}^M \sum_{s \in G_n(\hat{y}^{(i)})} C(s, \hat{y}^{(i)})} \quad (7)$$

**Definition 1.5** (Brevity penalty). Given a candidate corpus  $\hat{S} = (\hat{y}^{(1)}, \hat{y}^{(2)}, \dots, \hat{y}^{(M)})$  and corresponding reference corpi  $S = (S_1, S_2, \dots, S_M)$  where  $S_i = (y^{(i,1)}, y^{(i,2)}, \dots, y^{(i,N_i)})$ , the brevity penalty is defined as

$$BP(\hat{S}, S) = e^{-(r/c-1)^+} \quad (8)$$

where

$$c = \sum_{i=1}^M |\hat{y}^{(i)}| \quad (9)$$

is the length of the candidate corpus and

$$r = \sum_{i=1}^M \left| \operatorname{argmin}_{y \in S_i} (||y| - |\hat{y}^{(i)}||) \right| \quad (10)$$

is the effective reference corpus length.

**Definition 1.6** (Bilingual Evaluation Understudy (BLEU)). Given a candidate corpus  $\hat{S} = (\hat{y}^{(1)}, \hat{y}^{(2)}, \dots, \hat{y}^{(M)})$ , a corresponding reference corpi  $S = (S_1, S_2, \dots, S_M)$  where  $S_i = (y^{(i,1)}, y^{(i,2)}, \dots, y^{(i,N_i)})$ , and  $w \in \{[0, 1]^\infty : \sum_{i=1}^\infty w_i = 1\}$ , the BLEU score is defined as:

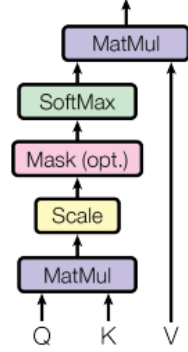
$$BLEU_w(\hat{S}, S) = BP(\hat{S}, S) \cdot \exp \left( \sum_{n=1}^\infty w_n \ln p_n(\hat{S}, S) \right) \quad (11)$$

**Definition 1.7** (ROUGE-n). Given a candidate sentence  $\hat{y}$  and a corresponding corpus of reference sentences  $S = \{y^{(1)}, y^{(2)}, \dots, y^{(N)}\}$ ,  $ROUGE - n$  is defined as

$$ROUGE - n(\hat{y}, S) = \sum_{i=1}^N \frac{\sum_{s \in G_n(y^{(i)})} \min(C(s, \hat{y}), C(s, y))}{\sum_{s \in G_n(y^{(i)})} C(s, y)} \quad (12)$$

Note that the definition of ROUGE metrics varies across sources. The definitions above align with [Sch17].

Scaled Dot-Product Attention



Multi-Head Attention

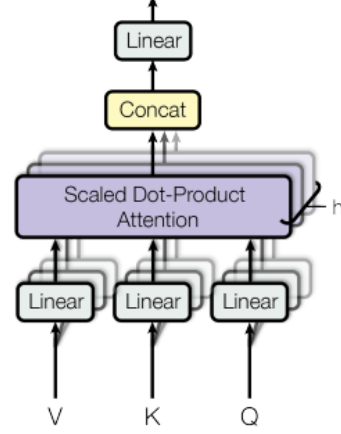


Figure 2: (left) Scaled Dot-Product Attention. (right) Multi-Head Attention consists of several attention layers running in parallel.

Figure 1: Attention Mechanism [VSP<sup>+</sup>17]

## 2 Attention

### 2.1 Dot-Product Attention

**Definition 2.1** (Dot-Product Attention). Given  $Q \in \mathbb{R}^{d_l \times d_k}$ ,  $K \in \mathbb{R}^{d_s \times d_k}$  and  $V \in \mathbb{R}^{d_s \times d_v}$

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V \in \mathbb{R}^{d_l \times d_v} \quad (13)$$

### 2.2 Multihead Attention

Given input matrices  $X_Q \in \mathbb{R}^{d_l \times d_g}$ ,  $X_K \in \mathbb{R}^{d_s \times d_w}$  and  $X_V \in \mathbb{R}^{d_s \times d_u}$  and weight matrices

$$\begin{aligned} \{W_Q^i \in \mathbb{R}^{d_g \times d_k}\}_{i=1}^h \\ \{W_K^i \in \mathbb{R}^{d_w \times d_k}\}_{i=1}^h \end{aligned} \quad (14)$$

$$\begin{aligned} \{W_V^i \in \mathbb{R}^{d_u \times d_v}\}_{i=1}^h \\ W_O \in \mathbb{R}^{hd_v \times d_o} \end{aligned} \quad (15)$$

we define

$$\begin{aligned} Q^i &= X_Q W_Q^i \in \mathbb{R}^{d_l \times d_k} \\ K^i &= X_K W_K^i \in \mathbb{R}^{d_s \times d_k} \\ V^i &= X_V W_V^i \in \mathbb{R}^{d_s \times d_v} \end{aligned} \tag{16}$$

$$A^i = \text{Attention}(Q^i, K^i, V^i) \in \mathbb{R}^{d_l \times d_v} \tag{17}$$

$$\text{Multihead}(X_Q, X_K, X_V) = \text{concat}(A_1, \dots, A_h) W_O \in \mathbb{R}^{d_l \times d_o} \tag{18}$$

### 2.3 Multihead Self-Attention

In Multihead Self-Attention,  $X_Q = X_K = X_V$  so that  $d_{\text{model}} := d_g = d_w = d_u$ ,  $L := d_l = d_s$  and  $\text{Multihead}(X_Q, X_K, X_V) \in \mathbb{R}^{L \times d_o}$ . If we also have  $d_o = d_{\text{model}}$  then we get the Multihead Self-Attention described in [VSP<sup>+</sup>17] Attention Is All You Need

### 2.4 Transformer Block

A single transformer block (see Figure 2) takes an input tensor  $X_0$  and then:

1. Applies a multi-head attention layer to  $X_0$  to produce  $X_1$
2. Adds  $X_0$  and  $X_1$  and applies normalization to produce  $X_2$
3. Applies a fully connected feed forward layer to  $X_2$  to produce  $X_3$
4. Adds  $X_2$  and  $X_3$  and applies normalization to produce the final output  $X_4$

## 3 Attention is All You Need [VSP<sup>+</sup>17]

Multi-head self-attention was introduced in the paper Attention is All You Need [VSP<sup>+</sup>17], which used the architecture in Figure 3 for a translation task.

Inputs to the model are Byte Pair Encoded (BPE, section 14) tokens, which are mapped to a learned embedding space. In order for the model to make use of the order of the input sequence, information on sequence order is injected by adding in sine and cosine function of different frequencies:

$$\begin{aligned} PE_{(pos, 2i)} &= \sin(pos/10000^{2i/d_{\text{model}}}) \\ PE_{(pos, 2i+1)} &= \cos(pos/10000^{2i/d_{\text{model}}}) \end{aligned} \tag{19}$$

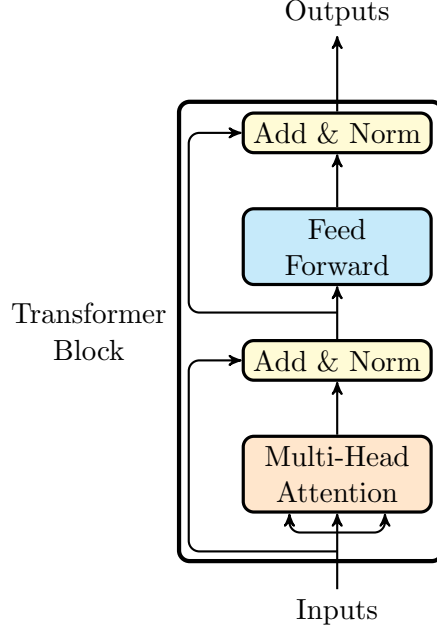


Figure 2: Single Transformer Block

## 4 Improving Language Understanding by Generative Pre-Training

[RNSS18] is the OpenAI article that introduced GPT (Generative Pre-Training). This framework uses unsupervised pre-training of a multi-layer transformer decoder with a masked language model objective followed by fine-tuning for specific tasks.

**Unsupervised pre-training** involves taking a corpus of tokens  $\mathcal{U} = \{u_1, u_2, \dots, u_n\}$  and training a model to maximize the likelihood objective:

$$L_1(\mathcal{U}) = \sum_i \log P(u_i | u_{i-k}, \dots, u_{i-1}; \Theta) \quad (20)$$

where  $k$  is the size of the context window,  $P$  is the conditional probability of the next token and  $\Theta$  are the parameters of the model.

[RNSS18] use a multi-layer transformer decoder for their model; a stack of multiple transformer blocks:

$$\begin{aligned} h_0 &= UW_e + W_p \\ h_l &= \text{transformer\_decoder}(h_{l-1}) \quad \forall i \in [1, \dots, n] \\ P(u) &= \text{softmax}(h_n W'_e) \end{aligned} \quad (21)$$

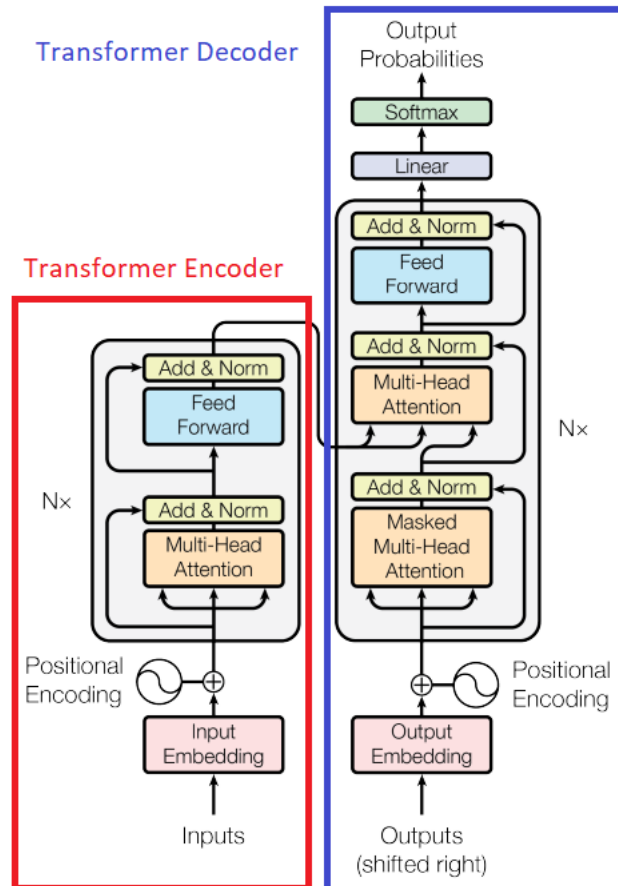


Figure 1: The Transformer - model architecture.

Figure 3: Transformer Architecture [VSP<sup>+</sup>17]

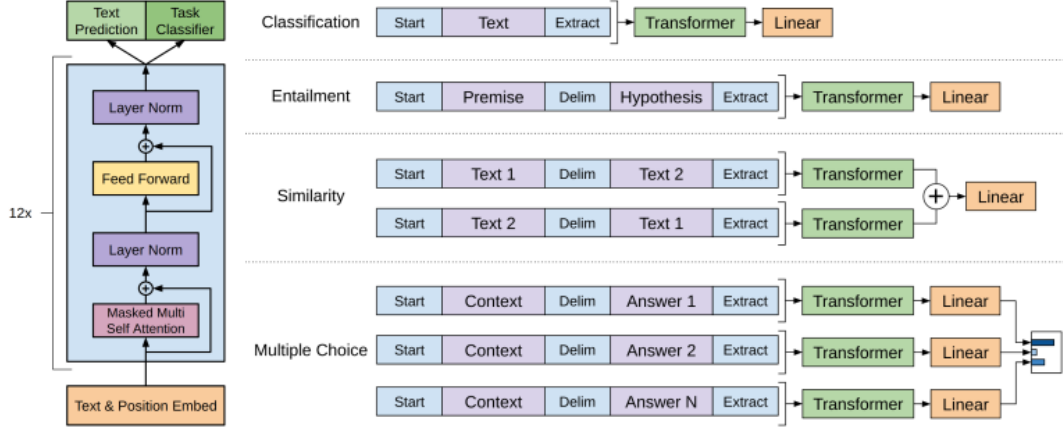


Figure 1: **(left)** Transformer architecture and training objectives used in this work. **(right)** Input transformations for fine-tuning on different tasks. We convert all structured inputs into token sequences to be processed by our pre-trained model, followed by a linear+softmax layer.

Figure 4: GPT Training Objectives [RNSS18]

where  $U = (u_{-k}, \dots, u_{-1})$  is the context vector of tokens,  $n$  is the number of layers,  $W_e$  is the token embedding matrix, and  $W_p$  is the position embedding matrix.

**Supervised fine-tuning** uses labeled training data to further optimize the model for certain tasks, including classification, semantic similarity, textual entailment and question answering. Given a labeled dataset  $\mathcal{C}$  of inputs  $\{x^1, \dots, x^m\}$  and targets  $y$ , the model is trained to maximize the likelihood objective:

$$L_2(\mathcal{C}) = \sum_{(x,y)} \log P(y|x^1, \dots, x^m) \quad (22)$$

where

$$P(y|x^1, \dots, x^m) = \text{softmax}(h_l^m W_y) \quad (23)$$

$h_l^m$  is the last output of the final transformer block and  $W_y$  are learned parameters.

During fine-tuning, the authors continued to use the masked language modeling to improve generalization and accelerate convergence. Hence, the final fine-tuning objective was

$$L_3(\mathcal{C}) = L_2\mathcal{C} + \lambda L_1(\mathcal{C}) \quad (24)$$

with a tuning parameter  $\lambda$  that trades off between the two objectives of performance on the fine-tuning task and generalization.

## **5 Language Models are Unsupervised Multitask Learners [RWC<sup>+</sup>18]**

[RWC<sup>+</sup>18] is the paper from OpenAI that introduced GPT-2. The principle finding of this paper was that by scaling up both model size from approximately 100M parameters to 1.5B parameters and using a significantly larger pre-training dataset (WebText - based on Common Crawl data with outbound Redit links), a Transformer-decoder type model could achieve state-of-the-art performance on many tasks without the need to fine-tune.

## **6 Language models are few-shot learners [BMR<sup>+</sup>20]**

[BMR<sup>+</sup>20] is the paper from OpenAI that introduced GPT-3. GPT-3 is again a transformer decoder-only model pre-trained on a curated version of the Common Crawl dataset. Again, OpenAI significantly scaled up the size of the model - the largest GPT-3 model has 175B parameters. Like GPT-2, GPT-3 was not fine-tuned for any specific task, but still demonstrated near-SOTA performance across a range of NLP tasks. GPT-3 performed better with infew-shot contexts; that is, when examples demonstrating desired output for the task was provided as part of the prompt.

## **7 Training language models to follow instructions with human feedback [OWJ<sup>+</sup>22]**

This is the paper from OpenAI that introduced Instruct GPT and the use of Reinforcement Learning from Human Feedback (RLHF). Instruct GPT uses a pre-trained GPT 3 model and fine-tunes it in three steps:

1. Given a set of prompts and desired responses, the model is fine-tuned on these examples in a supervised manor (SFT).
2. Given a prompt, several model outputs are sampled and a human labeler ranks them from best to worst. A dataset generated in this manor is then used to train a reward model (RM) (also a GPT 3 model).
3. Given a prompt, the model from 1 generates an output, the reward model calculates a reward and then the model is updated using reinforcement learning via proximal policy optimization.



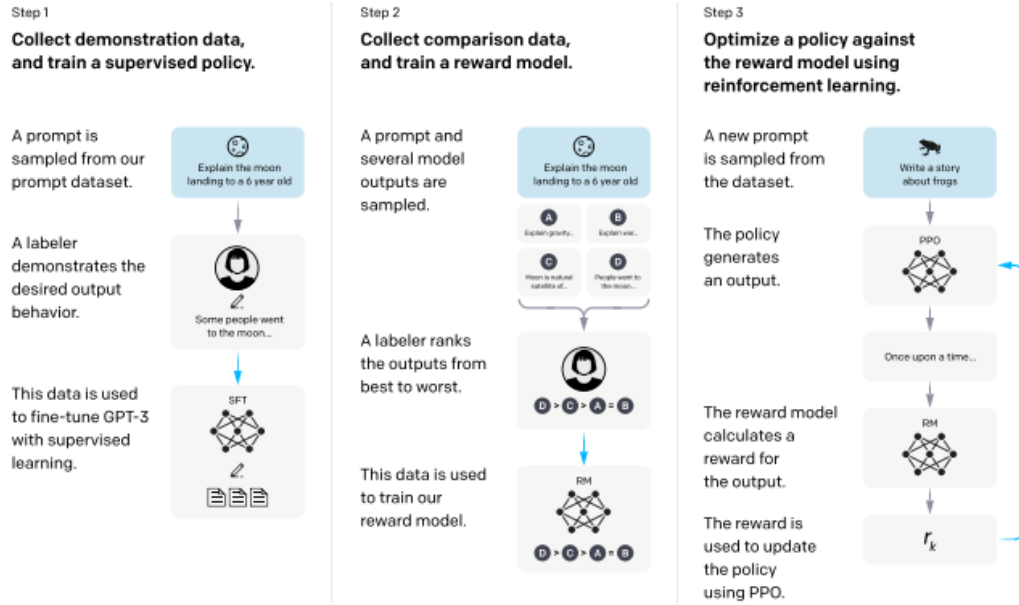


Figure 2: A diagram illustrating the three steps of our method: (1) supervised fine-tuning (SFT), (2) reward model (RM) training, and (3) reinforcement learning via proximal policy optimization (PPO) on this reward model. Blue arrows indicate that this data is used to train one of our models. In Step 2, boxes A-D are samples from our models that get ranked by labelers. See Section 3 for more details on our method.

Figure 5: Instruct GPT Fine-Tuning [OWJ<sup>+</sup>22]

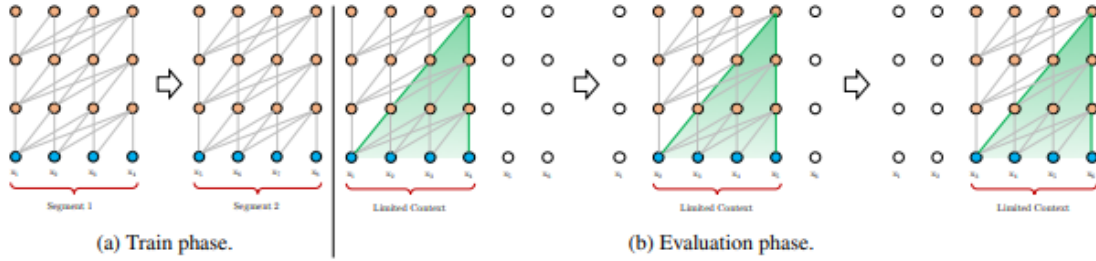


Figure 1: Illustration of the vanilla model with a segment length 4.

Figure 6: Transformer with Fixed Context Window [DYY<sup>+</sup>19]

## 8 BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

[DCLT19] was the first paper to introduce BERT models - transformer encoder only models that can be fine-tuned for a variety of tasks.

## 9 Scaling Transformers to Longer Sequences

One limitation of the original transformers is that their computational complexity grows quadratically in the length of the input sequence; that is, they have computational complexity of  $O(n^2)$ , where  $n$  is the sequence length. There have been several approaches developed to reduce this complexity.

### 9.1 Transformer-XL [DYY<sup>+</sup>19]

One crude option to reduce computational complexity during training is to split text into segments of length  $L$  and train a model on the individual segments, ignoring all contextual information from previous segments. This reduces complexity during training but causes contextual fragmentation as no information flows across segments. This vanilla model is shown in 6. At each time step during evaluation, the model processes a text segment of length  $L$ , the last output position is recorded and then the context window is shifted to the right by one step and the process repeated. By shifting only a single time step, each prediction is able to use the context of the last  $L$  positions, alleviating the contextual fragmentation in training, but reintroducing computational complexity.

In order to address the issue of contextual fragmentation, [DYY<sup>+</sup>19] introduce a recurrence mechanism. With two consecutive segments of length  $L$ ,  $s_\tau = [x_{\tau,1}, \dots, x_{\tau,L}]$  and  $s_{\tau+1} = [x_{\tau+1,1}, \dots, x_{\tau+1,L}]$ , let the  $n$ -th layer hidden state sequence produced by the  $\tau$ -th segment

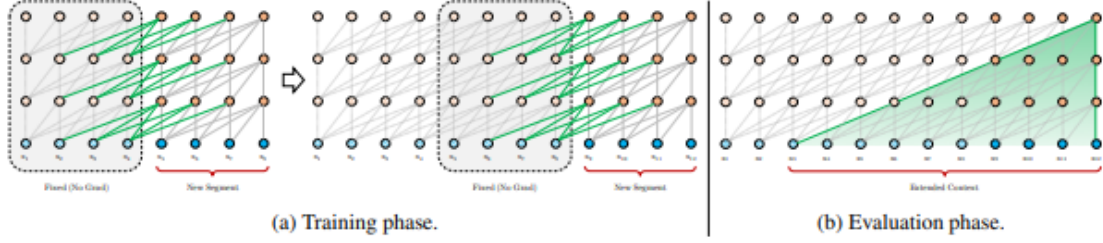


Figure 2: Illustration of the Transformer-XL model with a segment length 4.

Figure 7: TransformerXL [DYY<sup>+</sup>19]

$s_\tau$  be  $h_\tau^n \in \mathbb{R}^{L \times d}$ , where  $d$  is the hidden dimension. Then

$$\begin{aligned}\tilde{h}_{\tau+1}^{n-1} &= [SG(h_\tau^{n-1}) \circ h_{\tau+1}^{n-1}] \\ q_{\tau+1}^n, k_{\tau+1}^n, v_{\tau+1}^n &= h_{\tau+1}^{n-1} W'_q, \tilde{h}_{\tau+1}^{n-1} W'_k, \tilde{h}_{\tau+1}^{n-1} W'_v \\ h_{\tau+1}^n &= \text{transformer\_block}(q_{\tau+1}^n, k_{\tau+1}^n, v_{\tau+1}^n)\end{aligned}\tag{25}$$

where  $SG$  is the stop-gradient function and  $\circ$  denotes concatenation of vectors. Compared to the vanilla transformer, the keys and values in TransformerXL are constructed with an extended context, creating what is essentially a segment-level recurrence in hidden states.

## 9.2 Generating Long Sequences with Sparse Transformers

In [CGRS19], the authors scale transformers to longer sequences by factorizing the self-attention mechanism. Let  $S = \{S_1, \dots, S_L\}$ , where  $S_i \subseteq \{1, \dots, L\} \forall i \leq L$ . Then

$$\text{Attend}(X, S) = (a(X_{\cdot i}, S_i))_{i \in \{1, \dots, L\}}\tag{26}$$

$$a(X_{\cdot i}, S_i) = \text{softmax}\left(\frac{(W_q X_{\cdot i}) K'_{S_i}}{\sqrt{d}}\right) V_{S_i}\tag{27}$$

$$K_{S_i} = (W_k X_{\cdot j})_{j \in S_i} \quad V_{S_i} = (W_v X_{\cdot j})_{j \in S_i}\tag{28}$$

For full self-attention in an autoregressive model,  $S_i = \{j : j \leq i\}$  so that each element attends to all previous positions and its own position.

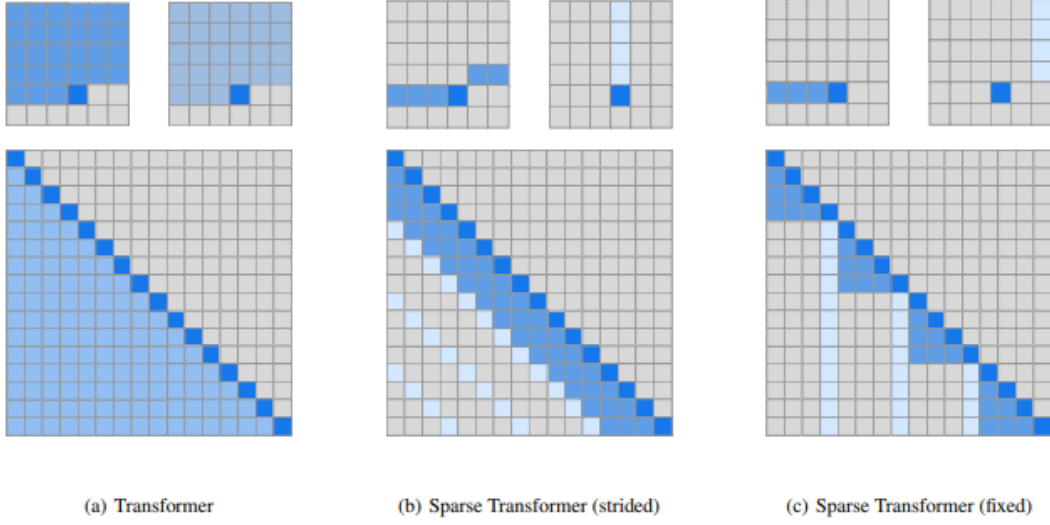


Figure 3. Two 2d factorized attention schemes we evaluated in comparison to the full attention of a standard Transformer (a). The top row indicates, for an example 6x6 image, which positions two attention heads receive as input when computing a given output. The bottom row shows the connectivity matrix (not to scale) between all such outputs (rows) and inputs (columns). Sparsity in the connectivity matrix can lead to significantly faster computation. In (b) and (c), full connectivity between elements is preserved when the two heads are computed sequentially. We tested whether such factorizations could match in performance the rich connectivity patterns of Figure 2.

Figure 8: Sparse Transformer Attention [CGRS19]

On the other hand, factorized self-attention uses  $p$  separate attention heads, where the  $m$ th head defines a subset of the indices  $A_i^{(m)} \subset \{j : j \leq i\}$ . In particular, the authors consider two factorizations:

- Strided Attention sets  $A_i^{(1)} = \{t, t+1, \dots, i\}$  for  $t = \max(0, i-l)$  and  $A_i^{(2)} = \{j : (i-j) \bmod l = 0\}$ .
- Fixed Attention sets  $A_i^{(1)} = \{j : \lfloor j/l \rfloor = \lfloor i/l \rfloor\}$  and  $A_i^{(2)} = \{j : j \bmod l \in \{t, t+1, \dots, l\}, t = l - c\}$ .

where  $l$  is the stride length and  $c$  is a hyperparameter. Strided attention performs best for data with structure that aligns with the stride, like images. Fixed attention performs well on data without a periodic structure, such as text.

### 9.3 Longformer [BPC20]

Longformer is a paper from AllenAI that introduced the Longformer model. Similar to sparse transformers, Longformer alters the attention pattern so that not all positions attend to all other positions. Longformer reduces the computational complexity  $O(n^2)$  of full

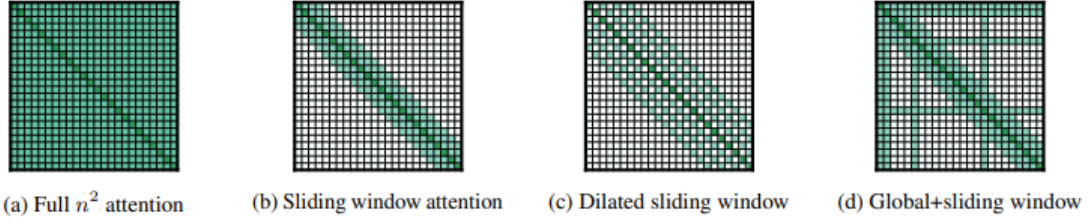


Figure 2: Comparing the full self-attention pattern and the configuration of attention patterns in our Longformer.

Figure 9: Longformer Attention [BPC20]

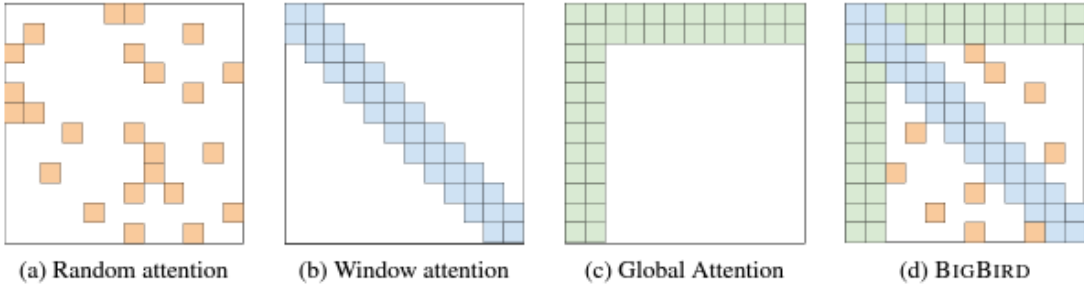


Figure 1: Building blocks of the attention mechanism used in BIGBIRD. White color indicates absence of attention. (a) random attention with  $r = 2$ , (b) sliding window attention with  $w = 3$  (c) global attention with  $g = 2$ . (d) the combined BIGBIRD model.

Figure 10: Big Bird Attention [ZGD<sup>+</sup>21]

attention to  $O(n \times w)$  by using a sliding window of size  $w$ . The paper also dilates the windows with gaps of size dilation  $d$  so that the final receptive field is  $l \times d \times w$ .

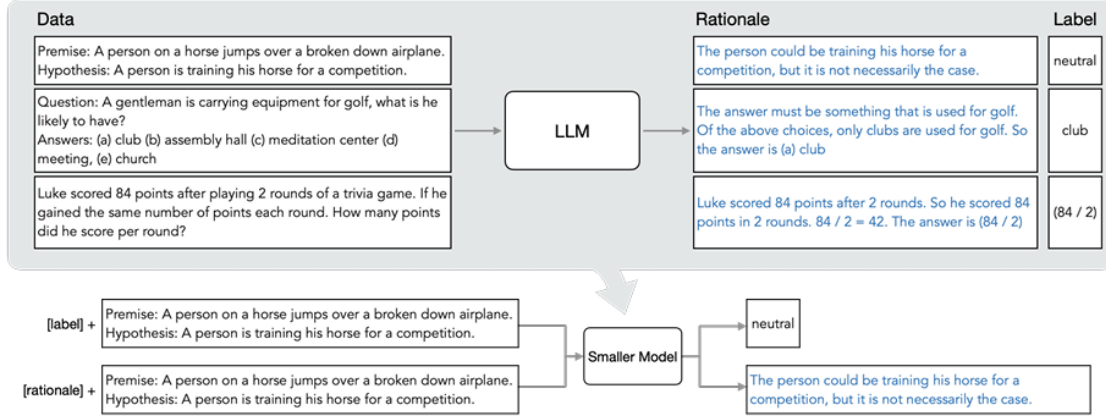


Figure 2: Overview on Distilling step-by-step. We first utilize CoT prompting to extract rationales from an LLM (Section 3.1). We then use the generated rationales to train small task-specific models within a multi-task learning framework where we prepend task prefixes to the input examples and train the model to output differently based on the given task prefix (Section 3.2).

Figure 11: Distilling Step-by-Step [HLY<sup>+</sup>23]

## 9.4 Big Bird: Transformers for Longer Sequences

# 10 Distilling

## 10.1 Distilling Step-by-Step! Outperforming Larger Language Models with Less Training Data and Smaller Model Sizes

# 11 Improving Fine-tuning

## 11.1 LoRA: Low-Rank Adaptation of Large Language Models [HSW<sup>+</sup>21b]

LoRA is a method for fine-tuning large, pre-trained transformers with limited fine-tuning data and/or compute. LoRA freezes pre-trained model weights and injects trainable rank decomposition matrices into each layer of the transformer architecture, significantly reducing the number of trainable parameters for downstream tasks.

Given a pre-trained weight matrix  $W_0 \in \mathbb{R}^{d \times k}$ , during fine-tuning we replace  $W_0$  with frozen( $W_0$ ) +  $BA$  where  $B \in \mathbb{R}^{d \times r}$  and  $A \in \mathbb{R}^{r \times k}$  are trainable parameters,  $\text{rank}(r) \ll \min(d, k)$ ,  $A$  has a random Gaussian initialization and  $B$  is initialized to zero.

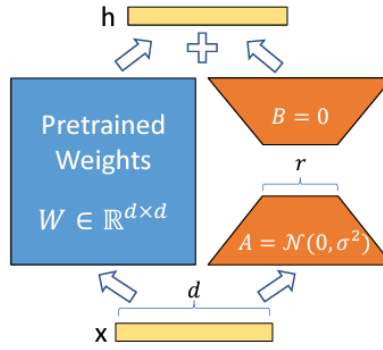


Figure 1: Our reparametrization. We only train  $A$  and  $B$ .

Figure 12: LoRA Fine-tuning Strategy [HSW<sup>+</sup>21a]

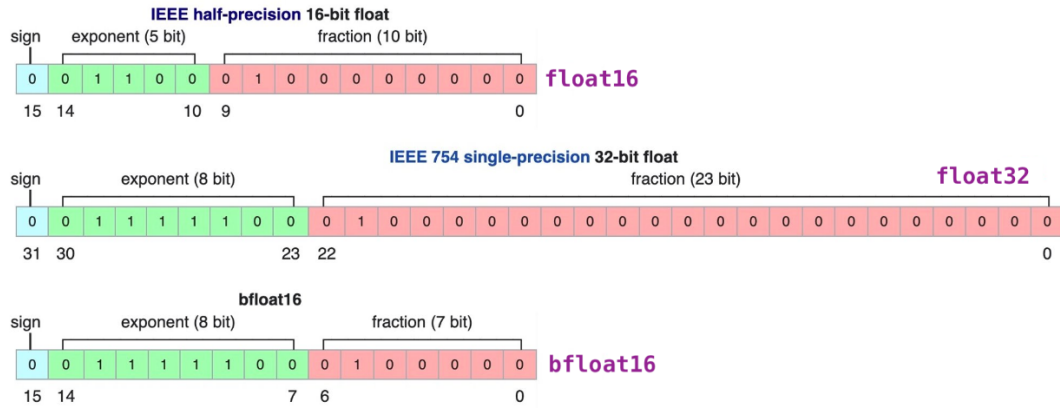


Figure 13: bfloat16 [bfl]

## 11.2 QLoRA

## 12 Relative Position Embeddings

Relative position embeddings were introduced by [SUV18].

Given relative position embedding matrix  $E^r \in \mathbb{R}^{L \times d_{model}}$  and matrix  $X \in \mathbb{R}^{L \times d_{model}}$

---

**Algorithm 1:** Relative Position Embedding of [HVV<sup>+</sup>18]

---

**Input:** Relative embedding matrix  $E^r \in \mathbb{R}^{L \times d_{model}}$

Matrix  $X \in \mathbb{R}^{L \times d_{model}}$

**Output:**  $D \in \mathbb{R}^{L \times L}$

```
1  $A \leftarrow X E^{rT} \in \mathbb{R}^{L \times L}$ 
2  $M \leftarrow \begin{cases} m_{i,j} = 1 & i \leq L, j \leq L, i \geq j \\ m_{i,j} = 0 & i \leq L, j \leq L, i < j \end{cases} \in \mathbb{R}^{L \times L}$ 
3  $A \leftarrow A \odot M$  // element-wise product
4  $B \leftarrow \begin{cases} b_{i,j} = 0 & j = 1, i \leq L \\ b_{i,j} = a_{i,j-1} & i \leq L, 1 < j \leq L+1 \end{cases} \in \mathbb{R}^{L \times L+1}$ 
5  $V \leftarrow \text{vec}(B^T)$ 
6  $C \leftarrow \{c_{ij} = V_{(i-1)L+j} \mid i \leq L, j \leq L\} \in \mathbb{R}^{L+1 \times L}$ 
7  $D \leftarrow \{d_{ij} = c_{i+1,j} \mid i \leq L, j \leq L\}$ 
8 return  $D$ 
```

---

## 13 Prompt Design and Tuning

Need to cover [LL21] and [LARC21]



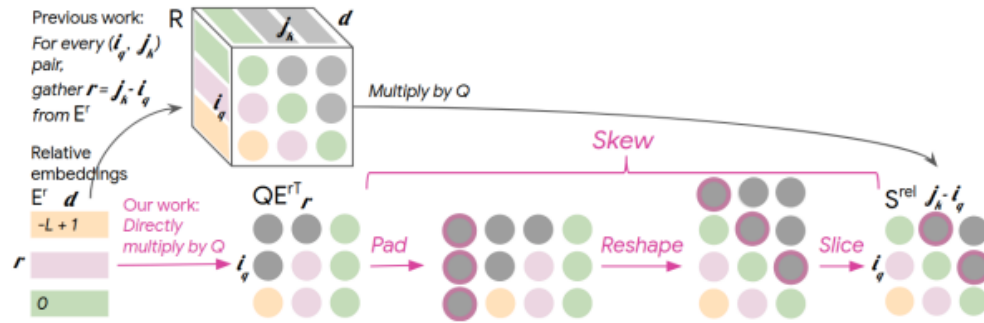


Figure 1: Relative global attention: the bottom row describes our memory-efficient “skewing” algorithm, which does not require instantiating  $R$  (top row, which is  $O(L^2D)$ ). Gray indicates masked or padded positions. Each color corresponds to a different relative distance.

Figure 14: Relative Global Attention [HVVU<sup>+</sup>18]

## 14 Byte-Pair Encoding

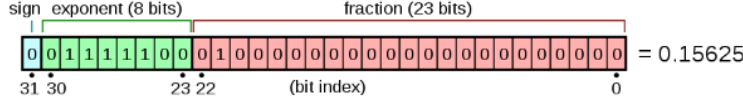
```
1 import re, collections
2
3 def get_stats(vocab):
4     pairs = collections.defaultdict(int)
5     for word, freq in vocab.items():
6         symbols = word.split()
7         for i in range(len(symbols)-1):
8             pairs[symbols[i],symbols[i+1]] += freq
9     return pairs
10
11 def merge_vocab(pair, v_in):
12     v_out = {}
13     bigram = re.escape(' '.join(pair))
14     p = re.compile(r'(?!\S)' + bigram + r'(?!\S)')
15     for word in v_in:
16         w_out = p.sub(' '.join(pair), word)
17         v_out[w_out] = v_in[word]
18     return v_out
19
20 vocab = {'l o w </w>' : 5, 'l o w e r </w>' : 2,
21         'n e w e s t </w>':6, 'w i d e s t </w>':3}
22
23 num_merges = 10
24
25 for i in range(num_merges):
26     pairs = get_stats(vocab)
27     best = max(pairs, key=pairs.get)
28     vocab = merge_vocab(best, vocab)
29     print(best)
```

Listing 1: Byte-Pair Encoding

## 15 Numeric Types and Precision

More papers to cover:

1. Roformer
2. FLAN T5 - Fine-tuning Language Net - instruction tuning
3. Roberta
4. Albert
5. Switch Transformer
6. GLaM



The real value assumed by a given 32-bit *binary32* data with a given *sign*, biased exponent  $e$  (the 8-bit unsigned integer), and a 23-bit *fraction* is

$$(-1)^{b_{31}} \times 2^{(b_{30}b_{29}\dots b_{23})_2 - 127} \times (1.b_{22}b_{21}\dots b_0)_2,$$

which yields

$$\text{value} = (-1)^{\text{sign}} \times 2^{(E-127)} \times \left(1 + \sum_{i=1}^{23} b_{23-i} 2^{-i}\right).$$

In this example:

- $\text{sign} = b_{31} = 0$ ,
- $(-1)^{\text{sign}} = (-1)^0 = +1 \in \{-1, +1\}$ ,
- $E = (b_{30}b_{29}\dots b_{23})_2 = \sum_{i=0}^7 b_{23+i} 2^{+i} = 124 \in \{1, \dots, (2^8 - 1) - 1\} = \{1, \dots, 254\}$ ,
- $2^{(E-127)} = 2^{124-127} = 2^{-3} \in \{2^{-126}, \dots, 2^{127}\}$ ,
- $1.b_{22}b_{21}\dots b_0 = 1 + \sum_{i=1}^{23} b_{23-i} 2^{-i} = 1 + 1 \cdot 2^{-2} = 1.25 \in \{1, 1 + 2^{-23}, \dots, 2 - 2^{-23}\} \subset [1; 2 - 2^{-23}] \subset [1; 2)$ .

thus:

- $\text{value} = (+1) \times 2^{-3} \times 1.25 = +0.15625$ .

Figure 15: float32 [flo]

7. QLoRA

8. Unigrams

9. SentencePiece

## References

- [bfl] To bfloat or not to bfloat? that is the question!  
<https://www.cerebras.net/machine-learning/to-bfloat-or-not-to-bfloat-that-is-the-question/>. Accessed: 2023-09-21.
- [BMR<sup>+</sup>20] Tom B Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *arXiv preprint arXiv:2005.14165*, 2020.
- [BPC20] Iz Beltagy, Matthew E. Peters, and Arman Cohan. Longformer: The long-document transformer, 2020.

- [CGRS19] Rewon Child, Scott Gray, Alec Radford, and Ilya Sutskever. Generating long sequences with sparse transformers, 2019.
- [DCLT19] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding, 2019.
- [DYY<sup>+</sup>19] Zihang Dai, Zhilin Yang, Yiming Yang, Jaime Carbonell, Quoc V. Le, and Ruslan Salakhutdinov. Transformer-xl: Attentive language models beyond a fixed-length context, 2019.
- [flo] Single-precision floating-point format. [https://en.wikipedia.org/wiki/Single-precision\\_floating-point\\_format](https://en.wikipedia.org/wiki/Single-precision_floating-point_format). Accessed: 2023-10-02.
- [HLY<sup>+</sup>23] Cheng-Yu Hsieh, Chun-Liang Li, Chih-Kuan Yeh, Hootan Nakhost, Yasuhisa Fujii, Alexander Ratner, Ranjay Krishna, Chen-Yu Lee, and Tomas Pfister. Distilling step-by-step! outperforming larger language models with less training data and smaller model sizes, 2023.
- [HSW<sup>+</sup>21a] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. *CoRR*, abs/2106.09685, 2021.
- [HSW<sup>+</sup>21b] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models, 2021.
- [HVU<sup>+</sup>18] Cheng-Zhi Anna Huang, Ashish Vaswani, Jakob Uszkoreit, Noam Shazeer, Ian Simon, Curtis Hawthorne, Andrew M. Dai, Matthew D. Hoffman, Monica Dinulescu, and Douglas Eck. Music transformer, 2018.
- [LARC21] Brian Lester, Rami Al-Rfou, and Noah Constant. The power of scale for parameter-efficient prompt tuning, 2021.
- [LL21] Xiang Lisa Li and Percy Liang. Prefix-tuning: Optimizing continuous prompts for generation, 2021.
- [OWJ<sup>+</sup>22] Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback, 2022.
- [RNSS18] Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. Improving language understanding by generative pre-training. 2018.

- [RWC<sup>+</sup>18] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. 2018.
- [Sch17] Natalie Schluter. The limits of automatic summarisation according to ROUGE. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 2, Short Papers*, pages 41–45, Valencia, Spain, April 2017. Association for Computational Linguistics.
- [SUV18] Peter Shaw, Jakob Uszkoreit, and Ashish Vaswani. Self-attention with relative position representations, 2018.
- [VSP<sup>+</sup>17] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need, 2017.
- [ZGD<sup>+</sup>21] Manzil Zaheer, Guru Guruganesh, Avinava Dubey, Joshua Ainslie, Chris Alberti, Santiago Ontanon, Philip Pham, Anirudh Ravula, Qifan Wang, Li Yang, and Amr Ahmed. Big bird: Transformers for longer sequences, 2021.